Sentiment Classification of Social Media

Project report submitted

in partial fulfilment of the requirement for the degree of

Masters in Computer Applications

by

Karan Solanki(2401560063)
Tanishq (2401560020)
Sakshi Yadav (2401560072)

Under the guidance of

Dr. Tannu Gupta

Department of Computer Science and Engineering

School of Engineering and Technology

K. R. Mangalam University, Gurugram – 122003

November – 2024

DECLARATION

We, the undersigned students of the **Master of Computer Applications** program at **K. R. Mangalam University, Gurugram**, hereby declare that the project report titled:

"Sentiment Classification of social media"

is a **bonafide record of our original work** carried out during the academic session **2024–2025**. This work has been completed in partial fulfilment of the requirements for the award of the **Master of Computer Applications (MCA)** degree under the guidance of **Dr. Tannu Gupta** our project supervisors.

We affirm that the work presented in this report is entirely our own, and has not been submitted elsewhere, in part or full, for the award of any other degree, diploma, or certificate from any other institution or university.

We fully understand that any violation of the above declaration may result in disciplinary action in accordance with the rules and regulations of the university.

Submitted by:

Student Name	**Roll No.**
Karan Solanki	2401560063
Tanishq	2401560020
Sakshi Yadav	2401560072

Place: K.R. Mangalam University

Date: 07 May, 2025

CERTIFICATE

It is certified that the work contained in the project report titled "Snake Game" by the student:

Name of the Student:	**Roll Number:**
Karan Solanki	2401560063
Tanishq	2401560020
Sakshi Yadav	2401560072

has been carried out under our/my supervision and that this work has not been submitted elsewhere.

Signature of Supervisors

Dr. Tannu Gupta
(External Suupervisor)

Date: –07.05.2024

Place: – K.R. Mangalam University

ACKNOWLEDGEMENT

We express our sincere gratitude to all those who supported us during the successful completion of our project titled:

"Sentiment Classification of social media"

We also extend our appreciation to our external mentor, **Dr. Tannu Gupta** for their insightful feedback and technical input.

Our heartfelt thanks to the **Department of Computer Science and Engineering**, **K. R. Mangalam University**, for providing the resources and academic environment that enabled us to complete this work.

We are equally grateful to our peers for their helpful discussions, and to our families for their constant support and motivation.

Place: – K.R. Mangalam University
Date: – 07.05.2025

Karan Solanki
Tanishq
Sakshi Yadav

ABSTRACT

In the digital age, social media platforms have become powerful tools for individuals and organizations to express opinions, share experiences, and influence public perception. Analyzing these vast amounts of user-generated content can provide valuable insights into public sentiment. This project aims to develop a sentiment classification system that automatically identifies the sentiment polarity—positive, negative, or neutral—of social media text using natural language processing (NLP) and machine learning techniques. The project includes data collection, text preprocessing, feature extraction using Bag-of-Words, and training a classifier (e.g., Naive Bayes) to predict sentiments. The system is evaluated using performance metrics such as accuracy and F1-score to ensure its reliability. The outcome demonstrates that machine learning can effectively be used to analyze public sentiment, making the solution useful for applications such as brand monitoring, political analysis, and crisis management. Future improvements may include deep learning models, real-time monitoring, and multilingual support.

INDEX

INTRODUCTION

In today's digitally connected world, social media platforms such as Twitter, Facebook, Instagram, and Reddit have become major channels through which people express their thoughts, opinions, emotions, and experiences. With billions of posts shared daily, these platforms serve as a vast, real-time repository of public sentiment. Analyzing this user-generated content provides an opportunity to extract meaningful insights into public opinion, market trends, and societal behavior—insights that are highly valuable for businesses, political campaigns, mental health organizations, and researchers alike.

However, manually analyzing such massive amounts of text data is neither practical nor efficient. This is where Sentiment Analysis, also known as opinion mining, comes into play. Sentiment analysis is a subfield of Natural Language Processing (NLP) that involves determining the emotional tone behind a body of text. The main objective of sentiment analysis is to classify the sentiment expressed in text as positive, negative, or neutral.

This project focuses specifically on sentiment classification of social media content using machine learning techniques. It aims to design a system that can automatically evaluate the emotional polarity of a given text post. The methodology involves several key stages:

Data Collection – Gathering labeled social media data (e.g., tweets or comments) from a public dataset or API.

Data Preprocessing – Cleaning and transforming raw text by removing noise such as hashtags, mentions, punctuation, stopwords, and converting all text to lowercase.

Feature Extraction – Converting processed text into numerical representations using methods like Bag-of-Words or TF-IDF (Term Frequency–Inverse Document Frequency).

Model Training – Using classification algorithms such as Naive Bayes, Logistic Regression, or Support Vector Machines to train a sentiment classifier on labeled data.

Evaluation – Measuring the model’s performance using standard metrics like accuracy, precision, recall, and F1-score.

Prediction – Applying the trained model to classify the sentiment of new, unseen social media posts.

The project not only demonstrates the effectiveness of NLP and machine learning in sentiment classification but also lays the foundation for future development in real-time opinion monitoring systems. This can benefit companies in managing customer relationships, political analysts in gauging public response, or even emergency services in detecting rising social tension during crises.

By building a system that automates sentiment classification, this project contributes toward creating intelligent applications that better understand human language and emotions—key steps toward more empathetic and responsive digital systems.

Key Features:

- Automated Sentiment Detection
  Classifies social media text (e.g., tweets, comments) into Positive, Negative, or Neutral sentiment categories using machine learning.

- Text Preprocessing Pipeline
  Implements NLP techniques such as lowercasing, punctuation removal, stopword removal, and tokenization to clean and normalize text.

- Customizable Dataset Input
  Supports structured input in .csv format, allowing users to train the model on various domains like product reviews, political tweets, or public feedback.

- Machine Learning Model Integration
  Uses Naive Bayes classifier for fast and effective sentiment prediction. Easily extendable to other algorithms like SVM, Logistic Regression, or deep learning models.

- Model Evaluation Tools
  Provides accuracy score, confusion matrix, and detailed classification report (precision, recall, F1-score) for performance analysis.

- Real-Time Sentiment Prediction
  Includes a function to input custom text and immediately receive sentiment classification results.

- Lightweight and Efficient
  Designed to be memory-efficient and fast, making it suitable for integration into lightweight apps or dashboards.

- Modular and Extensible Codebase
  Structured to allow easy upgrades, such as using TF-IDF, BERT, or integration with Flask/Streamlit for a web interface.

- Scalable for Big Data
  Can be scaled using tools like Spark NLP or integrated with APIs for live social media monitoring.

- Educational and Practical Use Cases
  Useful for students, researchers, and businesses wanting to monitor brand reputation or public mood.

OBEJECTIVE

- To collect and prepare a dataset of social media posts (e.g., tweets, comments, captions) labeled with sentiment categories.
- To preprocess the textual data by cleaning, tokenizing, removing stopwords, and normalizing the text for effective analysis.
- To extract relevant features from the text using methods like Bag of Words or TF-IDF for machine learning model input.
- To train and evaluate different machine learning algorithms (e.g., Naive Bayes, SVM, Logistic Regression) for classifying sentiment.
- To compare model performance using metrics such as accuracy, precision, recall, and F1-score to determine the most effective approach.
- To build a prediction function that allows sentiment classification of new, unseen social media text inputs.
- To provide insights or visualizations that highlight overall sentiment trends from the classified data.

PROBLEM FORMULATION AND OBJECTIVES

1. Problem Formulation

Social media has revolutionized the way people communicate and express their opinions. With millions of users posting content daily, platforms like Twitter, Facebook, Instagram, and Reddit have become rich sources of user sentiment and opinion. These expressions contain valuable information for organizations, political parties, service providers, researchers, and public health bodies. However, manually analyzing such massive, unstructured text data is impractical and time-consuming.

The challenge lies in converting this **unstructured natural language** into **structured, actionable insights**. Due to the informal and diverse nature of social media text—including abbreviations, slang, emojis, hashtags, and multilingual content—traditional text analysis methods often fall short.

Therefore, this project addresses the problem of **automatically detecting sentiment polarity** (positive, negative, or neutral) in user-generated social media posts. This is formally framed as a **supervised text classification problem**, where the goal is to train a machine learning model using labeled text data and use it to predict the sentiment of new, unseen posts.

Key challenges include:

- Handling noisy and informal language used in social media
- Reducing data dimensionality while preserving important linguistic features
- Choosing appropriate NLP techniques for preprocessing and feature extraction
- Selecting and tuning machine learning models for optimal classification performance
- Ensuring generalizability and robustness across diverse social contexts and domains

By solving this problem, the system can be used in a wide range of real-world applications, such as:

- Monitoring brand sentiment
- Understanding political opinions
- Detecting early signs of crises or unrest
- Analyzing customer reviews or public feedback

---

2. Objectives

The primary objective of this project is to design and develop a **sentiment classification system for social media** that can automatically analyze and categorize the emotional tone of textual content. The system should be able to process raw text data, transform it into usable input for machine learning models, and return sentiment predictions with high accuracy.

The project breaks this down into the following **specific, measurable objectives**:

Objective 1: Data Collection and Understanding

- To obtain a suitable dataset of social media posts labeled with sentiment classes (positive, negative, or neutral)
- To understand the structure, content, and quality of the dataset, and identify any issues such as class imbalance, missing data, or irrelevant entries

---

Objective 2: Text Preprocessing

- To implement a robust text preprocessing pipeline that cleans the raw text

- This includes:
    - Converting text to lowercase
    - Removing special characters, numbers, and URLs
    - Removing stopwords
    - Handling hashtags, mentions, and emojis
    - Tokenization and lemmatization or stemming

Objective 3: Feature Extraction

- To convert preprocessed text into numerical features suitable for machine learning
- To compare feature representation techniques such as:
    - Bag-of-Words (BoW)
    - Term Frequency–Inverse Document Frequency (TF-IDF)
    - Word Embeddings (e.g., Word2Vec, GloVe) for potential future improvements

---

Objective 4: Model Selection and Training

- To experiment with various supervised machine learning algorithms for text classification, including:
    - Naive Bayes
    - Logistic Regression
    - Support Vector Machine (SVM)
- To select the model that provides the best trade-off between accuracy, speed, and interpretability

---

Objective 5: Model Evaluation

- To evaluate model performance using standard metrics such as:
    - Accuracy
    - Precision
    - Recall
    - F1-Score
    - Confusion Matrix
- To apply cross-validation to assess generalization and reduce overfitting

---

Objective 6: System Deployment and Prediction

- To build a functional prediction module that takes custom user input and returns the predicted sentiment
- To consider integrating the system into a user interface or dashboard in future work

---

Objective 7: Exploratory Analysis and Insights

- To visualize sentiment trends using tools like matplotlib or seaborn
- To extract valuable insights, such as:
    - Most common words in each sentiment category
    - Volume of positive vs. negative sentiment over time

---

3. Alignment with Real-World Applications

This project is directly applicable in real-world domains such as:

- **Business intelligence** – Monitoring customer reviews to assess product reception
- **Politics** – Tracking shifts in public opinion during elections
- **Public health** – Detecting early signs of mental health concerns through negative or depressive sentiment
- **Marketing** – Optimizing campaigns based on real-time public reaction

The sentiment classification system developed here can serve as a foundational module in larger applications involving **social listening**, **trend forecasting**, and **automated content moderation**.

METHODOLOGY OF THE PROJECT

The methodology of the **Sentiment Classification of Social Media** project involves a step-by-step approach using techniques from Natural Language Processing (NLP) and supervised machine learning. The aim is to build a reliable model that can predict whether a given social media post conveys a **positive**, **negative**, or **neutral** sentiment.

The complete process can be broken down into the following stages:

---

1. Data Collection

The first step involves acquiring a dataset that contains social media posts along with sentiment labels. This can be done by:

- Downloading a public dataset (e.g., Twitter Sentiment140, Kaggle datasets)
- Collecting tweets using the **Twitter API** (with Tweepy or snscrape)
- Ensuring the dataset contains a balanced number of positive, negative, and neutral posts

Example data format:

pgsql

CopyEdit

```
| Tweet                              | Sentiment |
|------------------------------------|-----------|
| I love the new features!           | Positive  |
| This is so frustrating and broken. | Negative  |
| It’s okay, nothing special.        | Neutral   |
```

---

2. Data Preprocessing

Raw social media text is noisy and informal. Preprocessing cleans the text for machine learning by:

- **Lowercasing**: Converting all text to lowercase for uniformity
- **Removing punctuation**: Eliminates unnecessary symbols
- **Removing URLs, mentions (@user), hashtags (#)**: These add little value
- **Stopword removal**: Removes common words like "is", "the", etc.
- **Tokenization**: Splits text into words
- **Stemming/Lemmatization**: Reduces words to their base or root form

Example:

vbnet

CopyEdit

```
Original: “This new update is awesome! 😍 #loveit”
Processed: [‘new’, ‘update’, ‘awesome’, ‘love’]
```

---

3. Text Representation (Feature Extraction)

Since machine learning models work on numerical data, text must be converted into numerical features using methods like:

- **Bag of Words (BoW)**: Represents text as a vector of word counts
- **TF-IDF (Term Frequency-Inverse Document Frequency)**: Weighs word importance based on frequency and rarity
- **Word Embeddings** *(optional for future work)*: Dense vector representations (e.g., Word2Vec, GloVe)

---

4. Model Building (Machine Learning)

Several supervised learning algorithms can be used to classify sentiment based on the extracted features. In this project, we focus on:

- **Naive Bayes** – Fast, suitable for text classification
- **Logistic Regression** – Good for binary and multiclass classification
- **Support Vector Machine (SVM)** – Effective for high-dimensional data

Steps:

- Split the dataset into training and testing sets (e.g., 80% train, 20% test)
- Train the model on labeled data
- Validate and tune hyperparameters if needed

---

5. Model Evaluation

To assess the model's performance and ensure it generalizes well:

- Use metrics such as:
    - **Accuracy**
    - **Precision**
    - **Recall**
    - **F1-Score**
    - **Confusion Matrix**
- Perform **cross-validation** to reduce the risk of overfitting

Example confusion matrix:

markdown

CopyEdit

```
           Predicted
          | Pos | Neg | Neu
Actual | Pos   | 120 |  10 | 5
       | Neg   |  8  | 100 | 7
       | Neu   | 10  |  9  | 90
```

---

6. Sentiment Prediction

After training and testing, the final model can:

- Take new social media text as input
- Apply the same preprocessing and feature extraction
- Output predicted sentiment label (Positive / Negative / Neutral)

---

7. (Optional) Visualization

To enhance understanding and explainability:

- Use **word clouds** to visualize common words in each sentiment
- Use **bar plots or pie charts** to show distribution of sentiments
- Track **sentiment trends over time**

---

8. Tools and Libraries Used

- **Python 3.x**
- **Libraries**:
    - pandas, numpy – data handling
    - nltk, re, scikit-learn – NLP and machine learning
    - matplotlib, seaborn – visualization
    - Tweepy or snscrape – for live data (optional)

Summary Flowchart

nginx

CopyEdit

Raw Data → Preprocessing → Feature Extraction → Model Training → Evaluation → Prediction

Application

Applications of Sentiment Classification of Social Media
Sentiment classification is a powerful tool with widespread real-world applications across various industries. By analyzing the emotional tone behind user-generated content, organizations and individuals can gain insights, make informed decisions, and improve services or communication strategies. Below are some of the most impactful applications:

---

1. 🏢 Brand Monitoring and Reputation Management

Companies use sentiment analysis to track what customers are saying about their products or services on social media. This helps them:

- Detect negative feedback early and respond proactively
- Improve customer service
- Measure the effectiveness of marketing campaigns
- Understand overall brand perception

Example: A company tracks tweets mentioning its product. If there's a spike in negative sentiment, it investigates the issue and updates the product or sends apologies.

---

2. 🎯 Marketing and Campaign Optimization

Businesses analyze sentiment trends to gauge audience reactions to new launches, promotions, or advertisements. This helps them:

- Tailor marketing messages for different customer segments
- Identify which campaigns were successful
- Avoid strategies that provoke negative responses

---

3. 🗳️ Political Sentiment Analysis

Governments, political parties, and researchers use sentiment analysis to:

- Understand public opinion about policies or leaders
- Monitor changes in sentiment during election cycles
- Predict voting trends based on sentiment trends

Example: During an election, parties monitor Twitter to see how citizens react to campaign speeches or policy announcements.

---

4. 🧠 Mental Health Monitoring

Sentiment classification can be used to identify signs of depression, anxiety, or emotional distress in social media posts. Mental health organizations and researchers may:

- Detect patterns indicating a user's declining mental state
- Design early intervention programs
- Support at-risk individuals based on social media behavior

---

5. 🚨 Crisis and Disaster Response

During natural disasters or crises, people often express fear, confusion, or requests for help on social media. Sentiment classification helps emergency services and NGOs:

- Identify locations and people in distress
- Monitor public reaction to crisis response

- Understand the psychological impact of disasters

---

6. 💬 Customer Feedback Analysis

Organizations use sentiment analysis to process customer reviews or surveys and:

- Understand strengths and weaknesses in their services
- Prioritize feature requests or complaints
- Improve user experience and satisfaction

---

7. 📈 Stock Market and Financial Predictions

Financial institutions and investors use sentiment from social media to:

- Predict market movements
- Assess public confidence in a company
- Analyze news impact on stock prices

---

8. 🌍 Public Policy and Social Research

Researchers and policy makers use sentiment classification to:

- Study societal attitudes toward critical issues (e.g., climate change, vaccinations)
- Detect hate speech or toxic behavior
- Gauge support for laws or reforms

---

9. 🎮 Entertainment Industry

Studios and production companies analyze sentiment around shows, movies, and celebrities to:

- Predict box office success
- Tailor content to audience preferences
- Manage online fan engagement

---

10. 🧠 Artificial Intelligence Assistants

Virtual assistants (like Siri or Alexa) can be enhanced with sentiment detection capabilities to:

- Respond empathetically
- Adjust tone or suggestions based on user emotions

Challenges and Limitations

Challenges and Limitations

While sentiment classification of social media content offers immense potential, it is not without its challenges. Social media platforms present a unique set of difficulties due to the informal, dynamic, and diverse nature of their content. These limitations can impact the performance, accuracy, and generalizability of sentiment analysis models.

---

1. 🧾 Informal and Noisy Text

Social media posts often contain informal language, including:

- Slang and abbreviations (e.g., "LOL", "idk", "brb")
- Misspellings and typos
- Emojis, GIFs, hashtags, and mentions
- Short, context-less posts (e.g., “great!”)

These elements make it difficult for traditional NLP techniques to understand the true meaning of a message, and preprocessing steps may not handle them perfectly.

---

2. 😳 Ambiguity and Sarcasm

Understanding sentiment requires comprehension of **context, tone, and intent**—areas where machines struggle:

- Sarcastic sentences often appear positive but convey negative meaning (e.g., “Oh great, another bug 🙄”)
- Ambiguous statements can be interpreted differently depending on the user or situation

Handling sarcasm and ambiguity remains one of the hardest challenges in sentiment analysis.

---

3. 🌍 Multilingual and Code-Mixed Content

Users often write in **multiple languages** or switch between languages in the same sentence (e.g., English + Hindi). This poses a challenge for models that are trained on monolingual data.

Example:

“Yeh phone toh bilkul bakwaas hai, totally disappointed!”

Most standard models trained on English data fail to handle such code-mixed inputs effectively.

---

4. 📊 Imbalanced Datasets

Sentiment datasets often suffer from **class imbalance**, where one sentiment (usually neutral or positive) dominates the others. This skews the classifier toward the majority class and affects performance on minority sentiments.

---

5. 🏷️ Labeling Subjectivity

Sentiment is inherently **subjective**—what one person sees as neutral, another might see as negative. As a result, manually labeled datasets may contain inconsistencies and bias.

---

6. ⚙️ Feature Extraction Limitations

Simple feature extraction techniques like Bag-of-Words or TF-IDF:

- Ignore word order and context
- Cannot capture subtle linguistic patterns like negation (e.g., "not good" = bad)

While deep learning methods such as LSTM or transformers (like BERT) address these issues, they require more data and computational resources.

---

7. 🕒 Real-Time Sentiment Analysis Complexity

Deploying a sentiment classifier in a real-time environment (e.g., monitoring live tweets) involves:

- Handling large volumes of data quickly
- Processing messages with minimal delay
- Constant model updates to adapt to evolving language use

Achieving scalability and speed without sacrificing accuracy is a significant engineering challenge.

---

8. 🔐 Privacy and Ethical Concerns

Using social media data raises **privacy and ethical issues**, such as:

- Collecting personal or sensitive information without consent
- Potential misuse of analysis (e.g., manipulating opinions, surveillance)
- Biased models reflecting societal prejudices (gender, race, etc.)

Ethical guidelines and fair-use policies must be followed to prevent harm.

---

9. 📉 Generalization Issues

A model trained on one platform (e.g., Twitter) may not perform well on another (e.g., Reddit or Instagram) due to:

- Different writing styles and content types
- Varying character limits, context, and user behavior

This limits the generalizability of sentiment classifiers

CONCLUSION

Sentiment classification of social media has emerged as a powerful and necessary tool in the age of digital communication. With the ever-growing use of platforms like Twitter, Facebook, and Instagram, users express their opinions and emotions publicly, creating vast volumes of unstructured textual data. By leveraging Natural Language Processing (NLP) and machine learning techniques, this project successfully demonstrated how sentiment classification can be used to extract meaningful insights from such data.

The project followed a structured approach beginning with data collection and preprocessing, followed by feature extraction, model training, and performance evaluation. Using algorithms like Naive Bayes and Logistic Regression, the sentiment of social media posts was classified into positive, negative, or neutral categories with reasonable accuracy. The implementation highlighted the importance of clean data, proper feature engineering, and careful model selection to achieve effective sentiment prediction.

Through this study, it is evident that sentiment analysis can play a critical role in various sectors such as marketing, customer service, politics, and mental health. However, several challenges persist, including handling sarcasm, slang, multilingual data, and ethical concerns related to privacy.

Despite these limitations, the project successfully laid the groundwork for building intelligent systems that can understand public opinion at scale. It provides a foundation for future enhancements using advanced deep learning models, real-time data processing, and multilingual support.

References

Pang, B., & Lee, L. (2008). Opinion mining and sentiment analysis. Foundations and Trends® in Information Retrieval, 2(1–2), 1–135. https://doi.org/10.1561/1500000011

Liu, B. (2012). Sentiment analysis and opinion mining. Morgan & Claypool Publishers.

Go, A., Bhayani, R., & Huang, L. (2009). Twitter sentiment classification using distant supervision. CS224N Project Report, Stanford University. https://www.cs.stanford.edu/people/alecmgo/papers/TwitterDistantSupervision09.pdf

Bird, S., Klein, E., & Loper, E. (2009). Natural Language Processing with Python: Analyzing Text with the Natural Language Toolkit. O'Reilly Media.

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... & Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. Journal of Machine Learning Research, 12, 2825–2830.

Twitter Developer Platform. (2023). Twitter API Documentation. Retrieved from https://developer.twitter.com/en/docs

Kaggle. (2022). Sentiment140 Dataset with 1.6 Million Tweets. Retrieved from https://www.kaggle.com/kazanova/sentiment140

Jurafsky, D., & Martin, J. H. (2021). Speech and Language Processing (3rd ed. draft). Stanford University. https://web.stanford.edu/~jurafsky/slp3/

Mohammad, S. M., & Turney, P. D. (2013). NRC emotion lexicon. National Research Council Canada.

Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Proceedings of NAACL-HLT, 4171–4186.